

Principper for Samtidigshed og Styresystemer

Concurrency Problems

René Rydhof Hansen

29 APR 2019

Learning Goals: Last time

After last lecture you

- ... can define what a **race condition** is
- ... can explain how **mutual exclusion** can be used to avoid race conditions
- ... can explain strategies for **achieving and implementing** mutual exclusion
- ... can define **mutex** and **semaphore** and explain how they work and where they are useful
- ... can explain how to **synchronise** two (or more) threads and why it may be necessary

Learning Goals

After today's lecture, you

- ... can define and explain the concept of **deadlock**
- ... can define and explain **Coffman's conditions** for deadlock
- ... can explain and use **deadlock prevention strategies**:
 - prevention
 - avoidance
 - detect-and-recover
- ... can define and explain the following concepts
 - **livelock**
 - **priority inversion**
- ... can use the **Dining Philosophers** example to explain concurrency issues

Concurrency Problems

Concurrency... what's the problem?

Different problem categories

- Race Conditions
 - Atomicity violations
 - Order violations
- Deadlock

Real code, real bugs

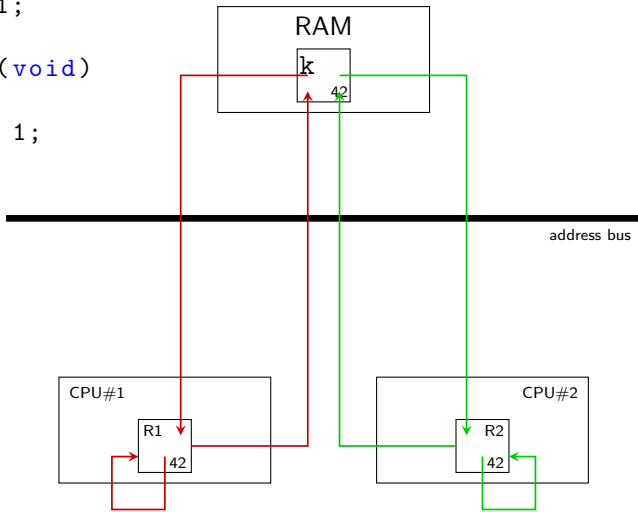
source: [OSTEP]

App	Non-deadlock	Deadlock
MySQL	14	9
Apache	13	4
Mozilla	41	16
OpenOffice	6	2
Total	74	31

Shared Memory Communication: Atomicity Violation

```
int k = 41;
```

```
void inck(void)
{
    k = k + 1;
}
```



Shared Memory Communication: Atomicity Violation

The Fix

```
int k = 41;

void inck(void)
{
    lock_mutex();
    k = k + 1;
    unlock_mutex();
}
```

Synchronisation: Order Violation

The Problem

```
float T;
```

```
void *thread0(void *tid)
{
    T = read_sensor();
}
```

```
void *thread1(void *tid)
{
    massive_computation(T);
    output();
}
```


Synchronisation: Order Violation

The Fix

```
float T;  
sem_t ready;  
sem_init(&ready,0,0);
```

```
void *thread0(void *tid)  
{  
    T = read_sensor();  
    sem_post(&ready);  
}
```

```
void *thread1(void *tid)  
{  
    sem_wait(&ready)();  
    massive_computation(T);  
    output();  
}
```

Deadlock: Basic Concepts

Deadlock

Definition (Deadlock)


A collection of threads P are in a **deadlock** state if every thread in P is waiting for an event that can only be generated by another thread in P .

Deadlock

Definition (Deadlock)

A collection of threads P are in a **deadlock** state if every thread in P is waiting for an event that can only be generated by another thread in P .

Dining Philosophers [Dijkstra]



- Five philosophers (think, eat)
- Must have (both) chopsticks before eating
- Must take **left** chopstick first

Deadlock

Definition (Deadlock)

A collection of threads P are in a **deadlock** state if every thread in P is waiting for an event that can only be generated by another thread in P .

Example (Bounded buffer [▶ details](#))

```
void *PROD() {  
    while(true) {  
        sem_wait(&mutex);  
        sem_wait(&free);  
        buffer[next_free] = data;  
        next_free =  
            (next_free + 1) % n;  
        sem_post(&used);  
        sem_post(&mutex);  
    }  
}
```

```
void *CONSUMER() {  
    while(true) {  
        sem_wait(&mutex);  
        sem_wait(&used);  
        data = buffer[next_used];  
        next_used =  
            (next_used + 1) % n;  
        sem_post(&free);  
        sem_post(&mutex);  
    }  
}
```

Deadlock

Definition (Deadlock)

A collection of threads P are in a **deadlock** state if every thread in P is waiting for an event that can only be generated by another thread in P .

Example (Bounded buffer ... with a problem! [▶ details](#))

```
void *PROD() {
    while(true) {
        sem_wait(&mutex);
        sem_wait(&free);
        buffer[next_free] = data;
        next_free =
            (next_free + 1) % n;
        sem_post(&used);
        sem_post(&mutex);
    }
}
```

```
void *CONSUMER() {
    while(true) {
        sem_wait(&mutex);
        sem_wait(&used);
        data = buffer[next_used];
        next_used =
            (next_used + 1) % n;
        sem_post(&free);
        sem_post(&mutex);
    }
}
```

Resource Allocation Graph

Definition (Resource Allocation Graph)


- Directed graph (V, E) with threads P and resources R as nodes:
 $V = P \cup R$
- An edge $(p, r) \in E$ if the thread $p \in P$ requests access to the resource $r \in R$
- An edge $(r, p) \in E$ if the thread $p \in P$ has access to resource $r \in R$

Deadlock


The threads P are in a deadlock state if and only if there is a **cycle** in the resource allocation graph^a

^aProvided there is only one instance of each resource

Resource Allocation Graph




Resource Allocation Graph



Control achieved
through **requests**

Resource Allocation Graph



Control achieved
through **requests**

Necessary conditions for deadlock

Coffman's Conditions (1971)

- **Mutual Exclusion**
 - Resources cannot be shared
- **No preemption**
 - Resources cannot be taken away from a thread
- **Hold-and-wait**
 - Threads may always try to take more resources
- **Circular Wait**
 - There is a circular dependency of resources

Necessary conditions for deadlock

Coffman's Conditions (1971)

- **Mutual Exclusion**
 - Resources cannot be shared
- **No preemption**
 - Resources cannot be taken away from a thread
- **Hold-and-wait**
 - Threads may always try to take more resources
- **Circular Wait**
 - There is a circular dependency of resources

Consequence

Deadlock can be **prevented** by breaking one or more of the conditions

Solution Strategies

Solution Strategies

1. Prevention (forebyggelse)

- Design that makes deadlock impossible (requires **proof**, e.g., through model-checking)
- Invalidate one or more of the conditions necessary for deadlock

2. Avoidance (undgåelse)

- Limited and controlled “lending” of resources
- Block threads with potentially dangerous allocation requests

3. Detection and Recovery (opdag og genopret)

- “Let’s see how bad it gets”

Solution: Deadlock prevention

Break one of Coffman's conditions

- **Avoid:** Mutual exclusion
 - Make resources shareable
 - Critical regions are **always** un-shareable resources
- **Allow:** Preemption
 - Allow pre-emption of resources (forcibly removing resources)
 - Only possible (safe) if state of resources can be re-created
- **Avoid:** Hold-and-wait
 - Require all resources to be allocated at once
 - Waste of resources
 - Potentially a long wait
- **Avoid:** Circular wait
 - Resources are allocated in a specific order
 - All resources must be known beforehand

Example: Bounded buffer and producer/consumer


[◀ return](#)

```
sem_t used, free;
int buffer[n],
    next_used = 0,
    next_free = 0;

int main() {
    sem_init(&used, 0, 0);
    sem_init(&free, 0, n);
    sem_init(&mutex, 0, 1);
}
```


```
void *PROD() {
    while(true) {
        sem_wait(&free);
        sem_wait(&mutex);
        buffer[next_free] = data;
        next_free = (next_free + 1) % n;
        sem_post(&mutex);
        sem_post(&used);
    }
}
```

```
void *CONSUMER() {
    while(true) {
        sem_wait(&used);
        sem_wait(&mutex);
        data = buffer[next_used];
        next_used = (next_used + 1) % n;
        sem_post(&mutex);
        sem_post(&free);
    }
}
```




Example: Producer/Consumer

- Mutual exclusion
- Preemption
- Allocate all resources at once
- Predefined allocation order




Example: Producer/Consumer

- Mutual exclusion
 - Cells can only be used by one thread at a time
 - Critical region must be protected
- Preemption
- Allocate all resources at once
- Predefined allocation order




Example: Producer/Consumer

- Mutual exclusion
 - Cells can only be used by one thread at a time
 - Critical region must be protected
- Preemption
 - Variables in a critical region may end up in an inconsistent state if involuntarily “kicked out” of critical region
- Allocate all resources at once
- Predefined allocation order




Example: Producer/Consumer

- Mutual exclusion
 - Cells can only be used by one thread at a time
 - Critical region must be protected
- Preemption
 - Variables in a critical region may end up in an inconsistent state if involuntarily “kicked out” of critical region
- Allocate all resources at once
 - Execute `sem_wait()` on both semaphores in an operation
- Predefined allocation order



Example: Producer/Consumer

- Mutual exclusion
 - Cells can only be used by one thread at a time
 - Critical region must be protected
- Preemption
 - Variables in a critical region may end up in an inconsistent state if involuntarily “kicked out” of critical region
- Allocate all resources at once
 - Execute `sem_wait()` on both semaphores in an operation
- Predefined allocation order
 - Exactly the chosen strategy for this example



Solution: Deadlock avoidance

Definition (Safe state)

A state in which there is **at least one** allocation sequence that allow all threads to finish without deadlock.

Definition (resource terminology)

Total amount of resources and total amount of **available** resources:

$$\vec{R} = (R_1, R_2, \dots, R_m) \qquad \vec{V} = (V_1, V_2, \dots, V_m)$$

Required resources (**claims**) and **allocated** resources:

$$\mathbf{C} = \begin{bmatrix} C_{11} & C_{12} & \cdots & C_{1m} \\ C_{21} & C_{22} & \cdots & C_{2m} \\ \vdots & \vdots & \vdots & \vdots \\ C_{n1} & C_{n2} & \cdots & C_{nm} \end{bmatrix} \qquad \mathbf{A} = \begin{bmatrix} A_{11} & A_{12} & \cdots & A_{1m} \\ A_{21} & A_{22} & \cdots & A_{2m} \\ \vdots & \vdots & \vdots & \vdots \\ A_{n1} & A_{n2} & \cdots & A_{nm} \end{bmatrix}$$

where C_{ij} (A_{ij}) requirement (current allocation) for process i resource j ;
with $R_j = V_j + \sum_{i=1}^n A_{ij}$ and $A_{ij} \leq C_{ij} \leq R_j$ for all i, j

Deadlock avoidance

Process Initiation Denial

With $n - 1$ processes running, only allow creation of process P_n if

$$\forall j: \quad R_j \geq C_{nj} + \sum_{i=1}^{n-1} C_{ij}$$

Pro's and Con's

- Effective (it works)
- Not efficient (too pessimistic)

Precondition

The maximum resource usage of **all** threads must be known **in advance**

Safe state

Definition

A system is in a **safe state** if there exists an ordering P_1, P_2, \dots, P_n of the system's threads such that

$$\forall i: \forall j: \quad C_{ij} - \sum_{i'=1}^i A_{i'j} \leq V_j$$

In words: a thread P_i can finish with the currently free resources and any resource(s) held by threads $P_{i'}$ for $i' \leq i$.

- In a safe state a system can **avoid** deadlock by letting processes run to termination starting with P_1
- Will an unsafe state always lead to deadlock?

Safe state

Definition

A system is in a **safe state** if there exists an ordering P_1, P_2, \dots, P_n of the system's threads such that

$$\forall i: \forall j: \quad C_{ij} - \sum_{i'=1}^i A_{i'j} \leq V_j$$

In words: a thread P_i can finish with the currently free resources and any resource(s) held by threads $P_{i'}$ for $i' \leq i$.

- In a safe state a system can **avoid** deadlock by letting processes run to termination starting with P_1
- Will an unsafe state always lead to deadlock? No... depends on scheduling etc.

Banker's Algorithm [Dijkstra 1965]

Definition

```
is_safe(C,A,V) {  
    Q = current set of processes  
    free = V  
    while (Q !=  $\emptyset$ )  
        remove some P[i] from Q such that  
            for all j:  $C[i,j] - A[i,j] \leq \text{free}[j]$   
            if no such p exists return FALSE  
            else  $\text{free}[j] = \text{free}[j] + A[i,j]$   
    return TRUE  
}
```

- Simulates allocation (of a **single** resource type) and checks if system is still in a safe state
- What if `is_safe` returns FALSE? **Unsafe state**

Banker's Algorithm: Properties and Limitations

- The maximum resource usage of **all** threads must be known **in advance**
 - Rarely realistic
 - Partial solution: use program analysis to approximate
- Only for **static number of processes**: rarely realistic
- Useful (maybe?) for specialised systems (RT, embedded)
- Presumes that processes with all necessary resources allocated **will return** resources (in finite time)

Example

	A	B	C	
$A_{0j} (C_{0j})$	2 (3)	1 (2)	0 (1)	P0: 1A + 1C? Not ok
$A_{1j} (C_{1j})$	1 (2)	1 (1)	1 (2)	P1: 1A + 1C? Ok
$V_j (R_j)$	1 (4)	0 (2)	2 (3)	

Solution: Detection and recovery

Detection and Recovery

- Assume deadlocks are rare
- Wait for deadlock to happen
- Re-establish normal (non-deadlocked) state

Preconditions for Detection and Recovery

- Deadlock state must be **detectable** (e.g., resource allocation graph)
- It must be possible to **define** and **re-establish** normal state

Definition (Ostrich algorithm (Strudsealgoritmen))

Bury head in sand and hope the problem goes away by itself, i.e., leave it to the user

Re-establishing normal state

Abort one of the deadlocked processes

- Acceptable if the process(es) can be restarted
- \LaTeX can be restarted, a network connection cannot always be “re-started” or re-established
- Which process should be terminated?

Check points

- Enables system to repeat aborted operations
- ... or to return to consistent state from before operation was aborted
- Example: transactional model for DB's
- Recovery points
- Problematic for processes with side-effects/external communication

Livelock, Starvation, and Priorities

Other “interesting” Concurrency Phenomena

- Livelock

- Deadlock-like state with no progress although with no cyclic wait
- Example: ethernet (exponential backoff protocol)

- Starvation

- Situation where a thread never acquires a resource because other processes always get it first
- Typical problem for systems with priorities

- Priorities

- What can have priorities: processes, threads, resources, ...
- Often useful/necessary with priorities
- Example: high priority for threads interacting with users (better response time)
- Problem: starvation, priority inversion

Priority Inversion

Priority Inversion

- When a low-priority thread blocks a high(er)-priority thread
- Happens when a higher-priority threads needs a resource held by a lower-priority thread

Example

Process L with **low** priority uses the printer; process H with **high** priority needs the printer.

- Result: H waits for L to finish. Problem?
- Alternative: Other processes with **medium** priority are scheduled instead of L which thus never finishes; de facto starvation of both H and L .

Example (Mars Pathfinder)

Low priority meteorological thread could not unlock mutex protecting communication bus needed for high priority thread.

Priority Inversion: Solutions

- **Interrupt disabling**

- Let interrupt masking be the only way to achieve mutex
- In reality only two priorities: preemptible and non-preemptible

- **Priority inheritance**

- A thread T inherits the priority of any higher priority threads blocked on resources held by T
- Dynamic calculation of ceiling and priority

- **Priority ceiling**

- Resources are assigned a **priority ceiling** corresponding to the highest priority thread that accesses the resource
- Threads using the resource are **temporarily** assigned the ceiling priority
- Requires code-/resource-analysis to determine ceiling

- **Immediate Ceiling Priority Protocol (ICPP)**

- Assign ceiling priorities to resources
- Temporarily raise the priority of a thread to the priority of accessed resources
- Threads can only request threads with higher priority

Properties of priority ceiling protocols

On a single processor

- A high-priority process is blocked **at most once** during its execution by lower-priority processes
- Deadlocks are **prevented**
- Transitive blocking is **prevented**
- Mutual exclusive access to resources is ensured by the protocol itself

Comparing OCPP versus ICPP

- Worst-case behaviour is **identical** (from a scheduling viewpoint)
- ICPP is easier to implement than the original (OCPP) as blocking relationships need not be monitored
- ICPP blocks prior to first execution: fewer context switches
- ICPP requires more priority movements as this happens with all resource usage
- OCPP changes priority only if an actual block has occurred

Alternative: Non-Blocking Data Structures

List insert (blocking)

```
void insert(int value) {  
    node_t *n = ...  
    n -> value = value;  
    mutex_lock();  
    n -> next = head;  
    head = n;  
    mutex_unlock();  
}
```

List insert (non-blocking)


```
void insert(int value) {  
    node_t *n = ...  
    n -> value = value;  
    do {  
        n -> next = head;  
    } while( CompareAndSwap(&head, n->next, n) == 0);  
}
```

Alternative: non-blocking data structures

- Solution at language-level
- Build non-blocking into data structure
- Example: unbounded non-blocking queue; uses compare-and-swap instruction directly


Dining Philosophers

Dining Philosophers [Dijkstra]




- If each philosopher picks up left chopstick, the result is deadlock
 - deny hold-and-wait: put chopstick down if the other is unavailable
 - this results in livelock: pick up, put down, pick up, put down
 - fastest philosopher will eventually break the livelock

Dining Philosophers [Dijkstra]




- Allow pre-emption: each philosopher grabs right-hand neighbour's chopstick
 - livelock again: get right chopstick, lose left chopstick, ...
 - strongest philosopher will eventually break the livelock

Dining Philosophers [Dijkstra]




- Allow sharing
 - how to share chopsticks?

Dining Philosophers [Dijkstra]




- Avoid circularities
 - add one extra chopstick so both neighbouring philosophers have one chopstick that isn't shared

Dining Philosophers [Dijkstra]



- Add a mutex to arbitrate resource acquisition
 - only the philosopher with the (one and only) soy sauce bottle can pick up chopsticks

Dining Philosophers [Dijkstra]



- Require both chopsticks are available before picking up either: Two philosophers can starve one seated between them

Summary

Concurrency — Important Concepts

- **Race condition**

- When the result of a computation depends on the **relative speed** of the individual threads
- In other words: the result depends on the actual **interleaving** of the threads
- Hard to debug

- **Critical region (critical section)**

- Program fragment vulnerable to race conditions
- Danish: “kritisk region”

Critical regions must be executed under mutual exclusion

- **Mutual exclusion (mutex)**

- When only one thread (among many) can access a given resource or execute specific part of the program-text
- Danish: “gensidig udelukkelse”

- **Atomic**

- Event, or sequence of events, that happen(s) **uninterruptedly**

Concurrency Summary

- The relative speed of threads is **unpredictable**
- Identify **shared resources**
- Identify **critical regions**
- Ensure **mutual exclusion** in critical regions
- Do **not** use scheduling for mutual exclusion
- Do **not** assume specific ordering on thread wake up
- Avoid **busy wait**
- Check if libraries are thread-safe before using them
- Check for: race conditions, deadlocks, livelocks, starvation, priority inversion, ...