

Principper for Samtidig og Styresystemer

Threads and Processes

René Rydhof Hansen

18 FEB 2019

Learning Goals

After today's lecture you

- ... can define and explain the concept of a **process**
- ... can explain what a **process image** is
- ... can explain what a **process control block**, what it is used for and why it is needed
- ... can explain, in general terms, how process creation, switching and termination works
- ... can define and discuss **process states**
- ... can define and explain the concept of a **thread**
- ... can discuss when, where, and how (multi-)threads are useful
- ... can explain what a **thread control block**, what it is used for and why it is needed
- ... can discuss **implementation strategies** for thread support and explain the associated **trade-offs**

Processes

What is a process?

Definition (Process)

- The environment a program is executed in
- A **running** program (with all concomitant data)
- A “virtual machine”
 - ... but with **limited** access to e.g. the processor and I/O devices
 - **Isolated** from other processes

Why processes?

- ...

What is an Operating System?

An application manager

- Primary task: to run programs
- Support execution of several programs (sequentially/concurrently)
- Maximising performance

A resource manager

- Primary resource: execution time (CPU)
- External devices: disks, printers, networks, ...
- Support development: provide libraries for common resources
- Support advanced features: sharing, IPC

An abstraction


- Every process has its own machine
- Support development: high-level APIs/libraries
- Uniform representation of resources

Running a Program


Running a Program (v1)

- ① Load code from secondary memory
 - How? ... later lecture!
- ② Allocate memory for the process
 - Assume simple fixed memory model (virtual... later!)
 - Code
 - Application memory
 - Heap
 - Stack (SP)
- ③ Start executing
 - PC points to next instruction
- ④ Continue until program terminates

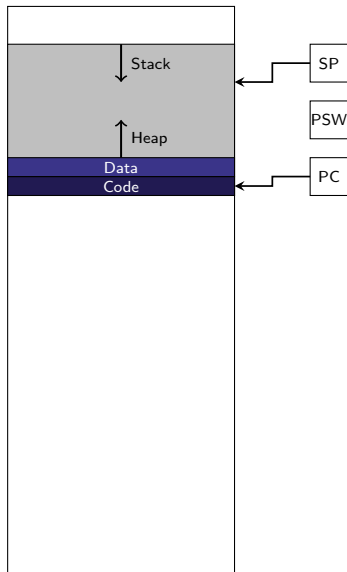
A Simple OS



A Simple OS



A Simple OS



A small problem

What happens at I/O?

- Wait for I/O to finish

A small problem

What happens at I/O?

- Wait a loooong time for I/O to finish
- Waste of time: CPU mostly idle
- Solution?

A small problem

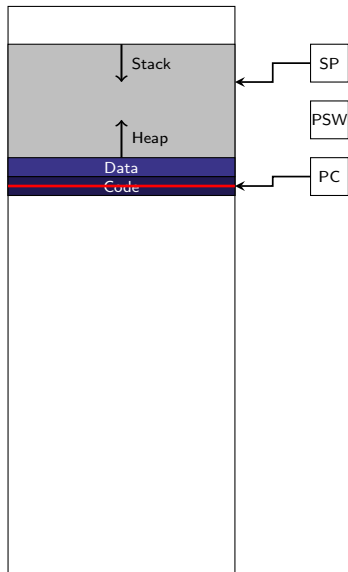
What happens at I/O?

- Wait a loooong time for I/O to finish
- Waste of time: CPU mostly idle
- Solution... switch to another program!

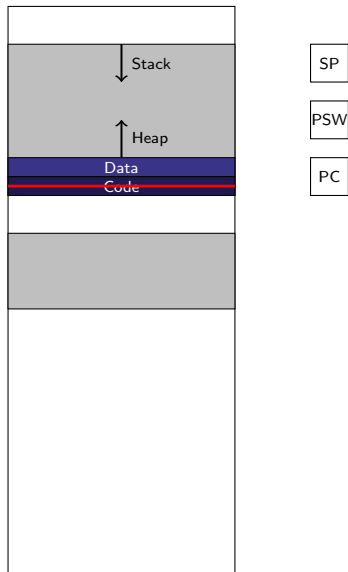
Running a Program (v2)

- ➊ Load code from secondary memory
- ➋ Allocate memory for the process
- ➌ Start executing
- ➍ Continue until program hits an I/O operation
- ➎ Switch to another program

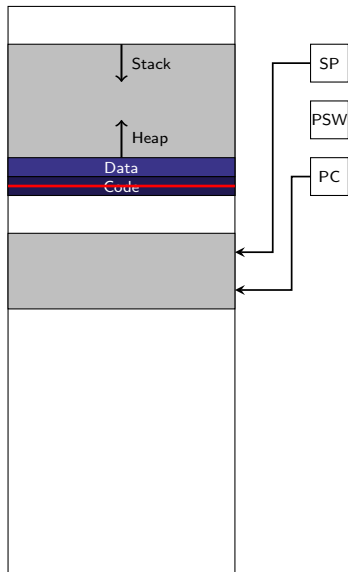
A Simple OS



A Simple OS



A Simple OS



The Lifecycle of a Process

When are processes created?

- Upon login
- Spawned by existing process
- New job started (by OS, user, ...)

When are processes terminated?

- Program completed
- Lack of resources
 - time (hard maximum limit)
 - memory
 - ...
- Error condition
 - invalid instruction
 - arithmetic error
 - ...
- Parent request/parent termination

How to switch back?

Easy!

- Save the relevant data for current process
- Re-instate relevant data for old process
- Problem?

Where to save this data?

- Reserved memory for each process: **Process Control Block**
- OS maintains **process table** linking PCBs

How to switch back?

Easy!

- Save the relevant data for current process
- Re-instate relevant data for old process
- Problem... relevant data?

Where to save this data?

- Reserved memory for each process: **Process Control Block**
- OS maintains **process table** linking PCBs

How to switch back?

Easy!

- Save the relevant data for current process
- Re-instate relevant data for old process
- Problem... relevant data?

Relevant Data

- Status registers: PC, SP, PSW,
- Data registers: EAX, ...
- Memory pointers (code, data, ...)
- I/O status

Where to save this data?


- Reserved memory for each process: **Process Control Block**
- OS maintains **process table** linking PCBs

Managing processes


Process Control Block (v1)

- Program counter
- Memory pointers
- Context data (registers)
- I/O status


A Simple OS




A Simple OS




A Simple OS




A Simple OS




A Simple OS



A Simple OS



A Simple OS



Running a Program (v3)

- ❶ Load code from secondary memory
- ❷ Allocate memory for the process
- ❸ Start executing
- ❹ Continue until program hits an I/O operation
- ❺ Switch to another program
 - Save the relevant data for current program
 - Pick another program
 - Re-instate new program's data
 - Execute

When to switch back?

- When the I/O operation has finished

When to switch back?

- When the I/O operation has finished
- ... and “the new” process has hit an I/O operation!
- In general: when OS regains control

When does the OS regain control?

- Interrupts
 - Clock interrupt
 - I/O interrupt
 - Memory fault
- Errors (traps)
- System calls

Important questions


- When does the OS execute?
- Is there an OS process?

How to choose the next process?

- Need to identify processes that are not waiting
 - Need to track the **state** of a program
 - Possible states: waiting, ready?


How to choose the next process?

- Need to identify processes that are not waiting
 - Need to track the **state** of a program
 - Possible states: waiting, ready?



How to choose the next process?

- Need to identify processes that are not waiting
 - Need to track the **state** of a program
 - Possible states: waiting, ready?
- Pick program from those that are ready
 - Which one?
 - Are some programs more important than others?
 - Need to track **priority** of a program
 - Enter PCBs into priority queue
 - **Scheduling**... more later!



Process Control Block (v2)

- Identifier
- State
- Priority
- Program counter
- Memory pointers
- Context data (registers)
- I/O status

Running a Program (v4)

- ① Load code from secondary memory
- ② Allocate memory for the process
- ③ Start executing
- ④ Continue until
 - ① Program executes blocking operation
 - ② Timeslice expires
- ⑤ Switch to another program
 - ① Save the relevant data for current program
 - ② Pick highest-priority program, ready to run
 - ③ Re-instate new program's data
 - ④ Execute

What is a process?

Definition (Process)

- The environment a program is executed in
- A **running** program (with all concomitant data)
- A “virtual machine”
 - ... but with **limited** access to e.g. the processor and I/O devices
 - **Isolated** from other processes

Process Image

Collection of process related data:

- Process Control Block
- Program (code)
- Stack
- Heap (user data)

What is a process?

Definition (Process (PSS))

- A **running** program (with all concomitant data)

Process Image

Collection of process related data:

- Process Control Block
- Program (code)
- Stack
- Heap (user data)

Threads

What is a process?

Definition (Process)

A **running** program (with all concomitant data)

Why processes?

- Unit of (resource) control
 - Devices (I/O, CPU), files (I/O), memory, locks, IPC, ...
- Unit of execution
 - Scheduling

Independent concerns?

What is a process?

Definition (Process)

A **running** program (with all concomitant data)

Why processes?

- Unit of (resource) control
 - Devices (I/O, CPU), files (I/O), memory, locks, IPC, ...
- Unit of execution
 - Scheduling

Independent concerns? ... to some degree: process as **context** for a thread

Definition (Thread)

The **executable** part of a process

Multithreading and Multiprogramming


Definition (Multiprogramming)

- Split of **tasks** that can be executed concurrently (in parallel?) on shared processors
- Implemented by rapid switching between tasks
- Common in modern OSs

Definition (Multithreading)

- Split of **processes** into several threads
- Always at least one thread pr. process (even in non-multithreaded systems)

Multithreading and Multiprogramming



Why threads?

Performance (perceived)

- Low-cost thread creation, switching, and termination
- Efficient inter-thread communication (shared memory)
- Can utilise multi-processor/multi-core platforms

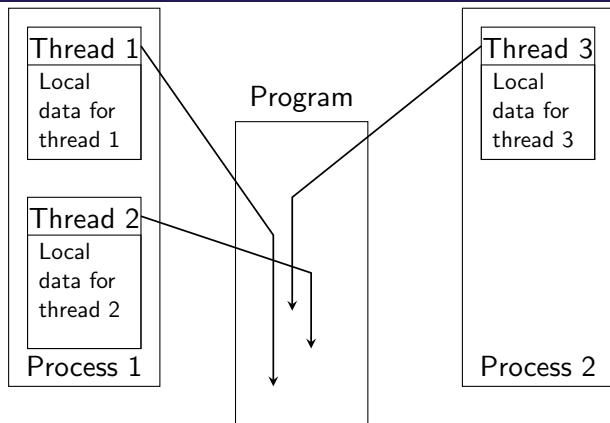
Better architecture/design

- Easy split into foreground and background work
- Asynchronous processing
- Modularity

Example (Responsive GUI)


- Some threads tasked with “main” job
- Some threads tasked with updating the GUI; better (perceived) responsiveness
- Better/easier design and implementation

Thread execution




- Two processes executing the **same** program(!)
- Program-text in **shared memory** (memory mapping)

Thread execution



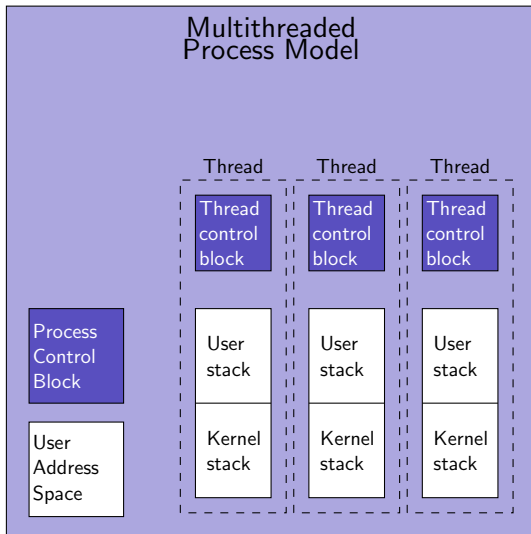

- Thread-local data?
- Process-local data?
- Shared data?

Thread execution



- Thread-local data? thread-ID, priority, stack
- Process-local data? address space, heap, open files, process-ID, parent, ownership, CPU reg. (copy)
- Shared data? Program-text

Process Control Block and Thread Control Block



Threads: Implementation Strategies

Implement in kernel-space (aka. obvious choice)

Implement in user-space

Hybrid: Kernel-support for user-space threads (aka. complex choice)

Threads: Implementation Strategies

Implement in kernel-space (aka. obvious choice)

- Modify/extend **kernel** to support threads
- Pro's: better handling of blocked threads

Implement in user-space

- **User-space library** for thread creation, switching, termination, ...
- Pro's: performance (kernel not involved), portable
- Con's: blocked threads

Hybrid: Kernel-support for user-space threads (aka. complex choice)

- Mainly in user-space
- Mapping user-space threads to (fewer) kernel-space threads
- Pro's: all of the above
- Con's: complex implementation, complex management

Sidetrack: Speedup of Multicore and Multithreading

Amdahl's Law: Exploiting Parallelism

Sidetrack: Speedup of Multicore and Multithreading

Amdahl's Law: Exploiting Parallelism

The potential performance speedup

$$\text{Speedup} = \frac{1}{(1 - f) + \frac{f}{N}}$$

where

- f is the “parallelisable fraction”
- N is the number of available cores/CPU's

Example

Assuming no overhead:

Speedup for $N = 8$ at $f = 1.0$: 8.0

Speedup for $N = 8$ at $f = 0.9$: 4.7 (with overhead: < 2)

Learning Goals

After today's lecture you

- ... can define and explain the concept of a **process**
- ... can explain what a **process image** is
- ... can explain what a **process control block**, what it is used for and why it is needed
- ... can explain, in general terms, how process creation, switching and termination works
- ... can define and discuss **process states**
- ... can define and explain the concept of a **thread**
- ... can discuss when, where, and how (multi-)threads are useful
- ... can explain what a **thread control block**, what it is used for and why it is needed
- ... can discuss **implementation strategies** for thread support and explain the associated **trade-offs**

Appendix: Process Control Block

- Identifier
- State
- Priority
- Program counter
- Memory pointers
- Context data (registers)
- I/O status
- Accounting info.