

# Principper for Samtidigshed og Styresystemer

## Scheduling

René Rydhof Hansen

25 FEB 2019

## Next time (04 MAR 2019): XV6 Hands-on Exercises

- No lecture
- TAs: 14:30–16:15 (approx)
- XV6 exercises (in group rooms)

# Learning Goals (last time)

After last lecture you

- ... can define and explain the concept of a **process**
- ... can explain what a **process image** is
- ... can explain what a **process control block**, what it is used for and why it is needed
- ... can explain, in general terms, how process creation, switching and termination works
- ... can define and discuss **process states**
- ... can define and explain the concept of a **thread**
- ... can discuss when, where, and how (multi-)threads are useful
- ... can explain what a **thread control block**, what it is used for and why it is needed
- ... can discuss **implementation strategies** for thread support and explain the associated **trade-offs**

# Learning Goals

After today's lecture you

- ... can explain the notion of **limited direct execution** and how it relates to scheduling
- ... will know the **simplified process model**
- ... will know and can explain important **metrics** for measuring a scheduling policy:
  - Fairness
  - Turnaround time
  - Response time
- ... can explain important **scheduling policies** and their pros and cons
  - FIFO
  - SJF
  - STCF
  - Round robin
  - MLFQ
  - Lottery scheduling

# Limited Direct Execution

# How to optimise CPU usage?

## Multitasking

- Processes **share** the CPU

## Problem(s)?

- **Efficient** execution of program
- Retaining **control** of running program
  - How to avoid dangerous/unwanted behaviour?
  - How to stop currently running program?

## The Answer: Limited Direct Execution

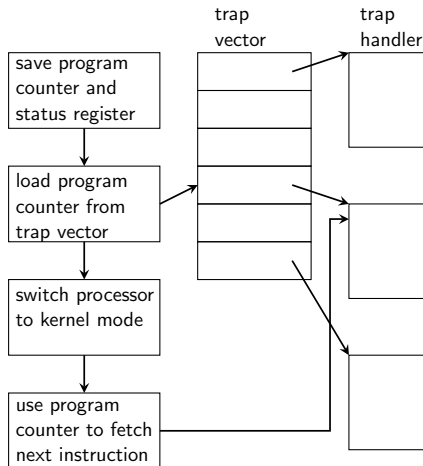
- User- vs- kernel-space (hw support)
- Timer interrupts (hw support)

# Recall from last time: Running a Program (v4)

- ➊ Load code from secondary memory
- ➋ Allocate memory for the process
- ➌ Start executing
- ➍ Continue until
  - ➊ Program executes blocking operation
  - ➋ Time slice expires
- ➎ Switch to another program
  - ➊ Save the relevant data for current program
  - ➋ Pick highest-priority program, ready to run
  - ➌ Re-instate new program's data
  - ➍ Execute

# Traps: Handling/implementation

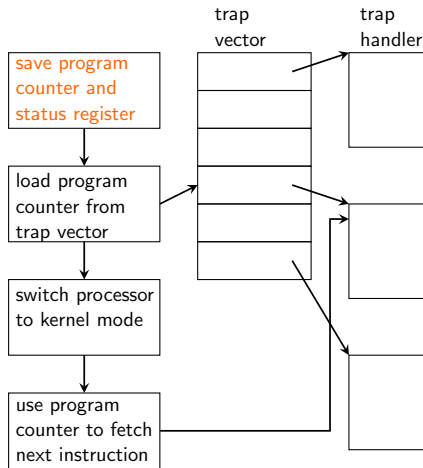
- **Trap handler** called
  - Program counter and status register saved on stack
  - Program counter set to address of exception handler
  - Kernel switches to kernel-mode
- Further information is saved
  - Current processor mode
  - Information about cause of exception
- **Trap vector**
  - Table with addresses of traps handlers
  - Initialised by operating system





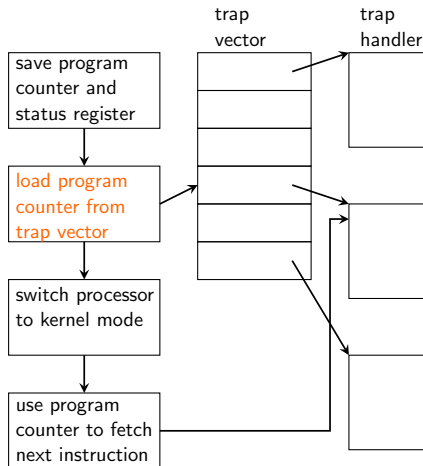
# Traps: Handling/implementation

- **Trap handler** called
  - Program counter and status register saved on stack
  - Program counter set to address of exception handler
  - Kernel switches to kernel-mode
- Further information is saved
  - Current processor mode
  - Information about cause of exception
- **Trap vector**
  - Table with addresses of traps handlers
  - Initialised by operating system



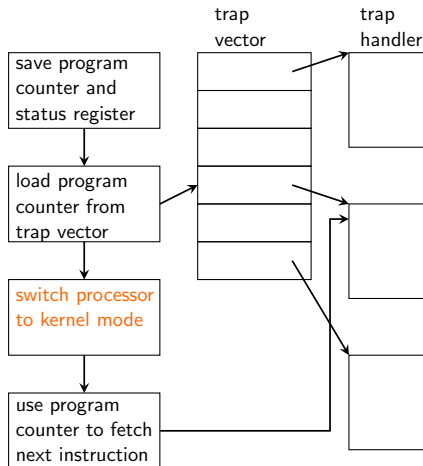
# Traps: Handling/implementation

- **Trap handler** called
  - Program counter and status register saved on stack
  - **Program counter set to address of exception handler**
  - Kernel switches to kernel-mode
- Further information is saved
  - Current processor mode
  - Information about cause of exception
- **Trap vector**
  - Table with addresses of traps handlers
  - Initialised by operating system



# Traps: Handling/implementation

- **Trap handler** called
  - Program counter and status register saved on stack
  - Program counter set to address of exception handler
  - **Kernel switches to kernel-mode**
- Further information is saved
  - Current processor mode
  - Information about cause of exception
- **Trap vector**
  - Table with addresses of traps handlers
  - Initialised by operating system



# Simple Scheduling

# Scheduling: Introduction

## Simplified process model

- ① Each job runs for the same amount of time.
- ② All jobs arrive at the same time.
- ③ Once started, each job runs to completion.
- ④ All jobs only use the CPU (i.e., they perform no I/O)
- ⑤ The run-time of each job is known.

## Metric: Turnaround Time

A metric for “how good” a scheduler is

$$T_{\text{turnaround}} = T_{\text{completion}} - T_{\text{arrival}}$$

## Terminology

For scheduling: job == process || job == thread

# Scheduling: FIFO

## FIFO

- Jobs are scheduled in the order they arrive

## Pros and Cons of FIFO Scheduling

- Easy to implement
- Little overhead
- Blissfully ignorant
- Fair
- “Convoy effect”

## Example (FIFO turnaround time)

Average turnaround for  $T_{exec}^i = 10$  ( $i = 1, 2, 3$ )?

# Relaxing Assumption #1

## Assumptions: Simplified process model

- ① Each job runs for the same amount of time.
- ② All jobs arrive at the same time.
- ③ Once started, each job runs to completion.
- ④ All jobs only use the CPU (i.e., they perform no I/O)
- ⑤ The run-time of each job is known.

# Relaxing Assumption #1

## Assumptions: Simplified process model

- 1 Each job runs for the same amount of time.
- 2 All jobs arrive at the same time.
- 3 Once started, each job runs to completion.
- 4 All jobs only use the CPU (i.e., they perform no I/O)
- 5 The run-time of each job is known.

## Example (FIFO turnaround time (re-visited))

Average turnaround for  $T_{exec}^1 = 10$ ,  $T_{exec}^2 = 10$ ,  $T_{exec}^3 = 100$ ? This is called the “convoy effect”



# Scheduling: Shortest Job First (SJF)

## Shortest Job First (SJF)

- Jobs are sorted **according to run-time length**
- Shortest (wrt. run-time) job is executed first

## Pros and Cons of SJF Scheduling

- Provably optimal
- Requires sorting of processes
- Fair?

# Scheduling: Shortest Job First (SJF)

## Shortest Job First (SJF)

- Jobs are sorted **according to run-time length**
- Shortest (wrt. run-time) job is executed first

## Pros and Cons of SJF Scheduling

- Provably optimal
- Requires sorting of processes
- Fair?

## Example (SJF turnaround time)

Average turnaround for  $T_{exec}^1 = 10$ ,  $T_{exec}^2 = 10$ ,  $T_{exec}^3 = 100$ ?

# Relaxing Assumption #2

## Assumptions: Simplified process model

- ① Each job runs for the same amount of time
- ② All jobs arrive at the same time
- ③ Once started, each job runs to completion
- ④ All jobs only use the CPU (i.e., they perform no I/O)
- ⑤ The run-time of each job is known

# Relaxing Assumption #2

## Assumptions: Simplified process model

- 1 Each job runs for the same amount of time
- 2 All jobs arrive at the same time
- 3 Once started, each job runs to completion
- 4 All jobs only use the CPU (i.e., they perform no I/O)
- 5 The run-time of each job is known

## Problem with SJF

- What if short(er) jobs arrive **later**?

## Example (SJF turnaround time (re-visited))

Average turnaround for  $T_{exec}^1 = 10$ ,  $T_{exec}^2 = 10$ ,  $T_{exec}^3 = 100$ , with  
 $T_{arrival}^1 \geq T_{arrival}^2 > T_{arrival}^3$ ?

# Relaxing Assumption #3

## Assumptions: Simplified process model

- 1 Each job runs for the same amount of time
- 2 All jobs arrive at the same time
- 3 Once started, each job runs to completion
- 4 All jobs only use the CPU (i.e., they perform no I/O)
- 5 The run-time of each job is known

# Relaxing Assumption #3

## Assumptions: Simplified process model

- 1 Each job runs for the same amount of time
- 2 All jobs arrive at the same time
- 3 Once started, each job runs to completion
- 4 All jobs only use the CPU (i.e., they perform no I/O)
- 5 The run-time of each job is known

## Preemption

- OS can **preempt** running process and (maybe) switch to new process
- Process-switch is expensive (overhead)

# Scheduling: Shortest Time-to-Completion First (STCF)

## Shortest Time-to-Completion First (STCF)

Whenever a (new) process enters the scheduler:

- Currently running process is **pre-empted**
- Jobs are sorted according to run-time **left**
- Job with **shortest remaining** run-time is allowed to run

## Pros and Cons of STCF

- Fair?
- Provably optimal

# Scheduling: Shortest Time-to-Completion First (STCF)

## Shortest Time-to-Completion First (STCF)

Whenever a (new) process enters the scheduler:

- Currently running process is **pre-empted**
- Jobs are sorted according to run-time **left**
- Job with **shortest remaining** run-time is allowed to run

## Pros and Cons of STCF

- Fair?
- Provably optimal
  - ... *if* execution time (job length) is known
  - ... *if* only turnaround is used
  - ... *if* no blocking operations are performed



# Response time

## Why turnaround time?

- Good for batch systems (optimises throughput)
- Good for interactive systems (latency)?

## Metric: Response time

$$T_{response} = T_{firstrun} - T_{arrival}$$

- **Note:** in RTS **response time** is defined more like turnaround-time

## Response time for FIFO, SJF, STFC

???

# Scheduling: Round Robin

## Round Robin

- Processes are entered into a **circular queue**
- Each process runs for a short while (**time slice**)
- OS switches to next process

## Pros and Cons of RR

- Fair?
- Poor **turnaround** time
- Good **response** time
- Tuning size of time slice (Warning: here be dragons!)

# Scheduling: Round Robin

## Round Robin

- Processes are entered into a **circular queue**
- Each process runs for a short while (**time slice**)
- OS switches to next process

## Pros and Cons of RR

- Fair?
- Poor **turnaround** time
- Good **response** time
- Tuning size of time slice (Warning: here be dragons!)

## Example (SJF vs RR response time)

Average response time for  $T_{exec}^i = 5$ , (for  $i = 1, 2, 3$ ) under SJF and RR (time slice = 1 time unit)

# Relaxing Assumptions #4 and #5

## Assumptions: Simplified process model

- ❶ Each job runs for the same amount of time
- ❷ All jobs arrive at the same time
- ❸ Once started, each job runs to completion
- ❹ All jobs only use the CPU (i.e., they perform no I/O)
- ❺ The run-time of each job is known

- Unrealistic(?): run-time known beforehand
  - (Hard-)RTS: necessary to know beforehand
- Unrealistic(?): No I/O

## Overlap

- Switch to ready process while waiting for I/O

# Relaxing Assumptions #4 and #5

## Assumptions: Simplified process model

- 1 Each job runs for the same amount of time
- 2 All jobs arrive at the same time
- 3 Once started, each job runs to completion
- 4 All jobs only use the CPU (i.e., they perform no I/O)
- 5 The run-time of each job is known

## Overlap

- Switch to ready process while waiting for I/O

# Not so simple scheduling

# Scheduling with priorities

## Scheduling goals

- Goal: optimise **turnaround** time
- Goal: minimise **response time**

## The Multi-Level Feedback Queue

- Assign **priorities** to processes
- Each priority level has its own process **queue**

## The Rules (MLFQ): How to schedule

Rule 1 If  $Pri(A) > Pri(B)$  run  $A$

Rule 2 If  $Pri(A) = Pri(B)$  use RR for  $A$  and  $B$

# Multi-Level Feedback Queue

## Changing priorities

- Claim: response time not/less important for CPU bound jobs
- Claim: jobs that often release CPU (before time): “interactive”

## The Rules (MLFQ): Changing priorities



# Multi-Level Feedback Queue

## Changing priorities

- Claim: response time not/less important for CPU bound jobs
- Claim: jobs that often release CPU (before time): “interactive”

## The Rules (MLFQ): Changing priorities

Rule 3 New processes are given **highest** priority

Rule 4a Priority of process using **all** its time slice is **reduced**

Rule 4b A process releasing the CPU before time, **stays** at the same priority

# Problems with the MLFQ

## Starvation

- What happens when many interactive (short) processes must be scheduled?

# Problems with the MLFQ

## Starvation

- What happens when many interactive (short) processes must be scheduled?
- Solution: the **priority boost**

## The Rules (MLFQ)

- Rule 1 If  $Pri(A) > Pri(B)$  run  $A$
- Rule 2 If  $Pri(A) = Pri(B)$  use RR for  $A$  and  $B$
- Rule 3 New processes are given **highest** priority
- Rule 4a Priority of process using **all** its time is **reduced**
- Rule 4b A process releasing the CPU before time, **stays** at the same priority
- Rule 5 “Boost” all processes at a fixed timed-interval

# Problems with the MLFQ

## Gaming the System

- What happens if a process uses **almost** all its time and then makes a system call? Repeatedly?

## The Rules (MLFQ)

Rule 1 If  $Pri(A) > Pri(B)$  run  $A$

Rule 2 If  $Pri(A) = Pri(B)$  use RR for  $A$  and  $B$

Rule 3 New processes are given **highest** priority

Rule 5 “Boost” all processes at a fixed timed-interval

# Problems with the MLFQ

## Gaming the System

- What happens if a process uses **almost** all its time and then makes a system call? Repeatedly?

## The Rules (MLFQ)

Rule 1 If  $Pri(A) > Pri(B)$  run  $A$

Rule 2 If  $Pri(A) = Pri(B)$  use RR for  $A$  and  $B$

Rule 3 New processes are given **highest** priority

Rule 4 Always **reduce** priority of a process when **alloted** time has been used

Rule 5 “Boost” all processes at a fixed timed-interval

# Proportional Share Scheduling

# Proportional Share Scheduling

## Fair share of CPU

- Optimise for CPU time (not response/turnaround time)

## The Lottery Scheduler

- Assign (a number of) tickets to each process
- Draw a (random) ticket
- Process with that ticket gets to run

## “Interesting” Features

- Probabilistic (no hard guarantees, only statistical)
- Trivial to implement
- Processes can **trade** tickets (why?)

## Stride Scheduling

- Each process is assigned a **stride** inversely proportional to number of tickets
- Each aprocess is assigned a **pass** number, incremented by the stride for every run
- The process with the lowest **pass** values runs



# Learning Goals

After today's lecture you

- ... can explain the notion of **limited direct execution** and how it relates to scheduling
- ... will know the **simplified process model**
- ... will know and can explain important metrics for measuring a scheduling policy:
  - Fairness
  - Turnaround time
  - Response time
- ... can explain important scheduling policies and their pros and cons
  - FIFO
  - SJF
  - STCF
  - Round robin
  - MLFQ
  - Lottery scheduling