

Principper for Samtidigshed og Styresystemer

Concurrency

René Rydhof Hansen

09 APR 2018

Next time...

XV6 hands-on exercise (v2): no lecture, no new exercises.

Learning Goals

After the last lecture, you are able to

- ... define and explain **paging** and how **paged memory** works
- ... perform simple **address translation** from paged (virtual) memory to physical memory
- ... explain how paged memory supports **shared memory**
- ... explain organisation of **page tables** (direct, two-level)
- ... define, explain, and discuss various **page replacement algorithms** and their pros and cons, including:
 - OPT (the optimal algorithm)
 - FIFO
 - Random
 - LRU
 - Clock

Learning Goals

After today's lecture you


- ... can define what a **race condition** is
- ... can explain how **mutual exclusion** can be used to avoid race conditions
- ... can explain strategies for **achieving and implementing** mutual exclusion
- ... can define **mutex**, **semaphore**, and **monitor** and explain how they work and where they are useful
- ... can explain how to **synchronise** two (or more) threads and why it may be necessary

Race Conditions


Shared Memory Communication: An Example

```
int k = 41;
```

```
void inck(void)
{
    k = k + 1;
}
```



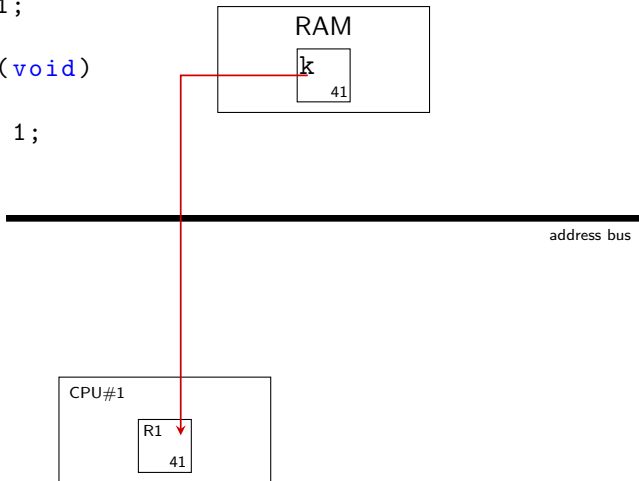
address bus



Shared Memory Communication: An Example

```
int k = 41;
```

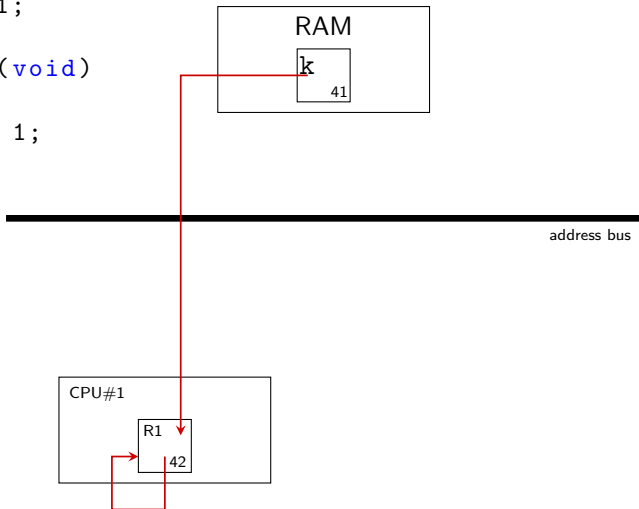
```
void inck(void)
{
    k = k + 1;
}
```



Shared Memory Communication: An Example

```
int k = 41;
```

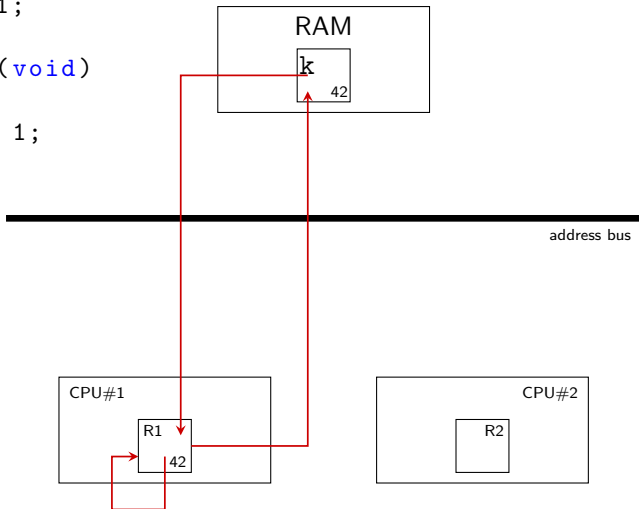
```
void inck(void)
{
    k = k + 1;
}
```



Shared Memory Communication: An Example

```
int k = 41;
```

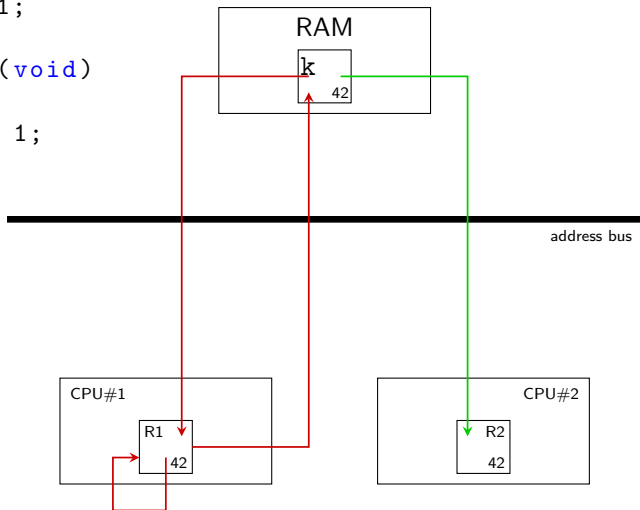
```
void inck(void)
{
    k = k + 1;
}
```



Shared Memory Communication: An Example

```
int k = 41;
```

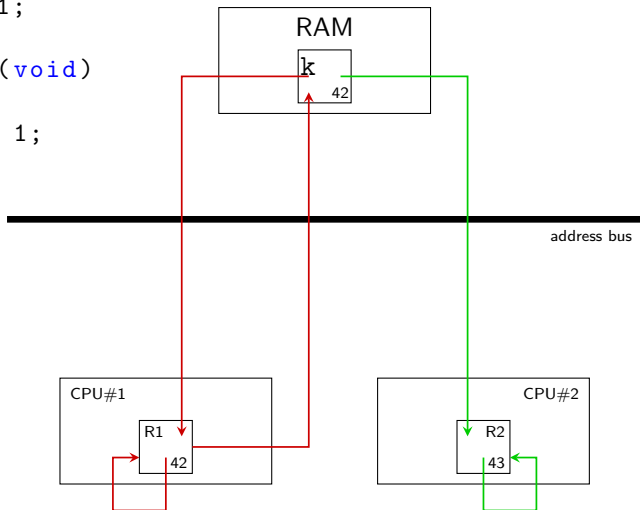
```
void inck(void)
{
    k = k + 1;
}
```



Shared Memory Communication: An Example

```
int k = 41;
```

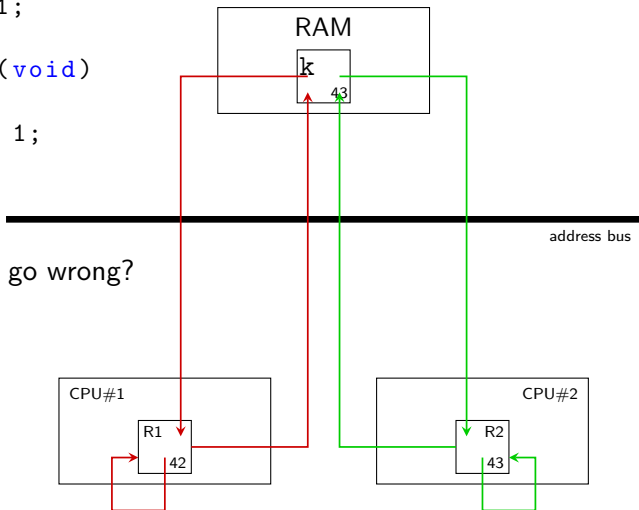
```
void inck(void)
{
    k = k + 1;
}
```



Shared Memory Communication: An Example

```
int k = 41;
```

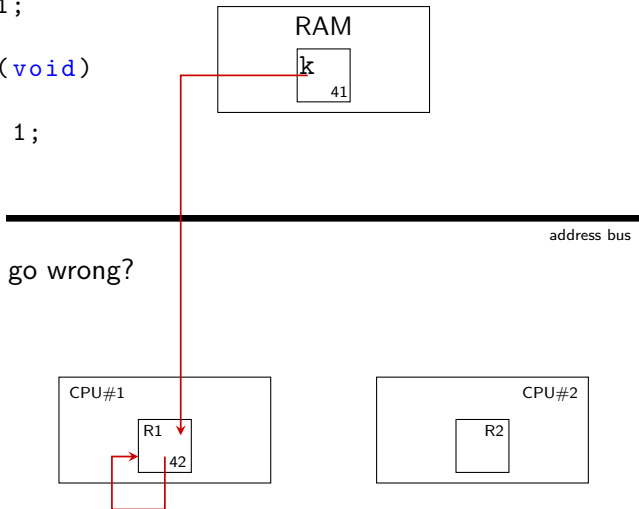
```
void inck(void)
{
    k = k + 1;
}
```



Shared Memory Communication: An Example

```
int k = 41;
```

```
void inck(void)
{
    k = k + 1;
}
```

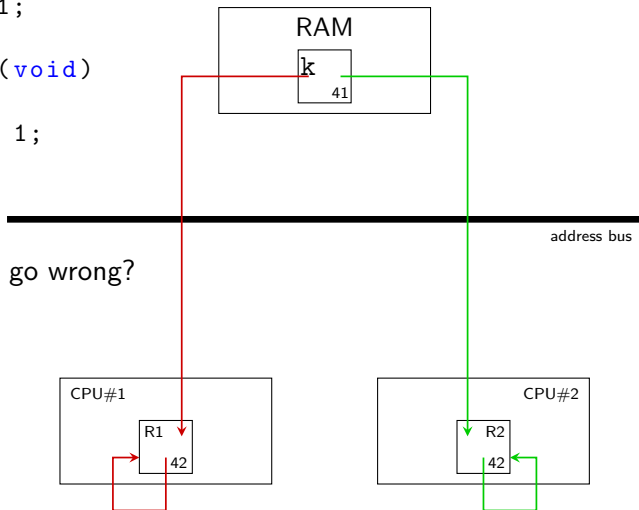


What could go wrong?

Shared Memory Communication: An Example

```
int k = 41;
```

```
void inck(void)
{
    k = k + 1;
}
```

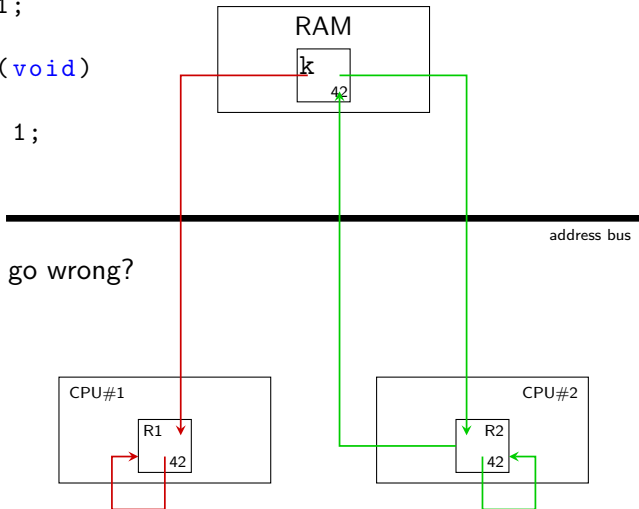


What could go wrong?

Shared Memory Communication: An Example

```
int k = 41;
```

```
void inck(void)
{
    k = k + 1;
}
```

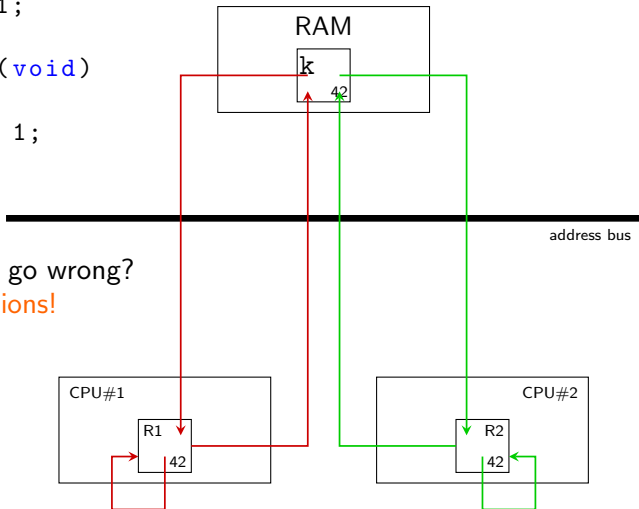


What could go wrong?

Shared Memory Communication: An Example

```
int k = 41;
```

```
void inck(void)
{
    k = k + 1;
}
```



What could go wrong?

Race conditions!

Concurrency — Important Concepts

- **Race condition**

- When the result of a computation depends on the **relative speed** of the individual threads
- In other words: the result depends on the actual **interleaving** of the threads
- Hard to debug

- **Critical region (critical section)**

- Program fragment vulnerable to race conditions
- Danish: “kritisk region”

Critical regions must be executed under mutual exclusion

- **Mutual exclusion (mutex)**

- When only one thread (among many) can access a given resource or execute specific part of the program-text
- Danish: “gensidig udelukkelse”

- **Atomic**

- Event, or sequence of events, that happen(s) **uninterruptedly**

Ensuring Mutual Exclusion

How to ensure mutual exclusion?

Locks!

- Ensure only **one** thread has access to critical region
- “Locking” critical regions

Evaluation Criteria for Locks

- Should guarantee **mutual exclusion** (i.e., it should work)
- Should be **fair**
- Should be **efficient**

How to ensure mutual exclusion?

Solution: Disabling Interrupts

- Disable interrupts before entering critical region
- Effective... but dangerous
- Works for **tiny** and **predictable** critical regions only
- Mostly relevant for **very low level** code, i.e., kernel code
- Difficult to incorporate in schedulability analysis, RTSs
- What about interrupts on multi-processor systems?

Code: Mutual exclusion by interrupt disabling

```
void *thread(void *tid) {  
    /* ... */  
    cli();  
    /* critical region */  
    sti();  
    /* ... */  
}
```

How to ensure mutual exclusion?

Solution: Software locks

- Use “flag” to indicate entry into critical region
 - ① Check if lock is open/closed, acquire lock
 - ② Access resources (protected by the lock)
 - ③ Release lock

Code: Programming with locks

```
lock_t  lock;

void  *thread(void *tid) {
    /* ... */

    acquire_lock(&lock);
    /* critical region */
    release_lock(&lock);

    /* ... */
}
```

Implementing locks

Lock variables

```
bool flag[2] = {false, false};

void *thread0(void *tid)    void *thread1(void *tid)
{
    while(flag[1]) { }      {
    flag[0] = true;          while(flag[0]) { }
    /* critical region */    flag[1] = true;
    flag[0] = false;         /* critical region */
    }                        flag[1] = false;
                            }
}
```

Implementing locks

Lock variables

```
bool flag[2] = {false, false};  
void *thread0(void *tid)    void *thread1(void *tid)  
{                            {  
    while(flag[1]) { }      while(flag[0]) { }  
    flag[0] = true;          flag[1] = true;  
    /* critical region */    /* critical region */  
    flag[0] = false;        flag[1] = false;  
}
```

The “algorithm” is **wrong**!

Tricky to implement

Implementation of locks susceptible to race conditions: when the while-loop ends, the “competing” process can take over again

Implementing locks

A working lock implementation: Dekker's Algorithm (1965)

```
bool flag[2] = {false, false};
```

```
int turn = 0;
```

```
void *thread0(void *) {  
    flag[0] = true;  
    while flag[1] {  
        if(turn == 1) {  
            flag[0] = false;  
            while(turn == 1) { }  
            flag[0] = true;  
        }  
    }  
}
```

```
/* critical region */
```

```
turn = 1;
```

```
flag[0] = false;
```

```
}
```

```
void *thread1(void *) {  
    flag[1] = true;  
    while flag[0] {  
        if(turn == 0) {  
            flag[1] = false;  
            while(turn == 0) { }  
            flag[1] = true;  
        }  
    }  
}
```

```
/* critical region */
```

```
turn = 0;
```

```
flag[1] = false;
```

```
}
```


Implementing locks (Dekker's Algorithm)

The Good

- It works!
- Efficient when blocking is not needed

The Bad

- Not very efficient in general
- A compiler may allocate some variables in registers (thus not shared between threads)
- Difficult (impossible?) to scale to more than two processes

Implementing locks: a better implementation

Peterson's mutex-algorithm (1981)

```
bool flag[2] = {false, false};
int  turn;

void *thread0(void *) {
    flag[0] = true;
    turn = 1;
    while(flag[1] &&
          turn == 1) { }
    /* critical region */
    flag[0] = false;
}

void *thread1(void *) {
    flag[1] = true;
    turn = 0;
    while(flag[0] &&
          turn == 0) { }
    /* critical region */
    flag[1] = false;
}
```

The Good and the Bad

- It works!
- Scales to more processes (... but still generally inefficient)
- Guarantees **fairness**
- Vulnerable to compilers optimising memory access

Implementing locks

The Ugly

Memory consistency on modern hardware?

How to ensure mutual exclusion?

Hardware supported locks

- Similar to software locks, but exploit special hardware
- Atomic instructions

Code: Programming with locks

```
int    lock;  
  
void   *thread(void *tid) {  
    /* ... */  
  
    while(locked(&lock) == 1) {}  
    /* critical region */  
    lock = 0;  
  
    /* ... */  
}
```

Hardware supported locks: test-and-set

The 'test-and-set' instruction

- Return **old** value and set to 1
- Executed **atomically**

Code: Informal semantics of test-and-set

```
int test_and_set(int *old_ptr) {  
    int old = *old_ptr;  
    *old_ptr = 1;  
    return old;  
}
```

Hardware supported locks: test-and-set

Code: Mutual exclusion using test-and-set

```
int lock = 0;

void *thread0(void *) {
    /* ... */
    while( t_a_s(&lock) == 1)
    {
        /* do nothing */
    }
    /* critical region */
    lock = 0;
    /* ... */
}

void *thread1(void *) {
    /* ... */
    while( t_a_s(&lock) == 1)
    {
        /* do nothing */
    }
    /* critical region */
    lock = 0;
    /* ... */
}
```

Hardware supported locks: xchg

The 'xchg' (exchange) instruction

- Swap arguments **atomically**

Code: Informal semantics of xchg

```
int xchg(int *old_ptr, int new) {  
    int old = *old_ptr;  
    *old_ptr = new;  
    return old;  
}
```

Hardware supported locks: xchg

Code: Mutual exclusion using xchg (from XV6)

```
void acquire(struct spinlock *lk) {
    pushcli(); // disable interrupts to avoid deadlock.
    if(holding(lk))
        panic("acquire");

    // The xchg is atomic.
    // It also serializes...
    while(xchg(&lk->locked, 1) != 0)
        ;

    // Record info about lock acquisition for debugging.
    lk->cpu = cpu;
    getcallerpcs(&lk, lk->pcs);
}
```


Hardware supported locks: compare-and-swap

The 'compare-and-swap' instruction

- Compare old value to test value, update if equal and **return old value**
- Executed **atomically**

Code: Informal semantics of `compare_and_swap`

```
int compare_and_swap(int *ptr, int expected, int new) {  
    int actual = *ptr;  
    if (actual == expected)  
        *ptr = new;  
    return actual;  
}
```

Hardware supported locks: compare-and-swap

Code: Mutual exclusion using compare-and-swap

```
int lock = 0;

void *thread0(void *) {
    /* ... */
    while( c_a_s(&lock,0,1)
           == 1)
    {
        /* do nothing */
    }
    /* critical region */
    lock = 0;
    /* ... */
}

void *thread1(void *) {
    /* ... */
    while( c_a_s(&lock,0,1)
           == 1)
    {
        /* do nothing */
    }
    /* critical region */
    lock = 0;
    /* ... */
}
```

Hardware supported locks

The 'fetch-and-add' instruction

- Atomically increment a variable and return old value

Code: Informal semantics for `fetch_and_add`

```
int fetch_and_add(int *ptr) {  
    int old = *ptr;  
    *ptr = old + 1;  
    return old;  
}
```

Hardware supported locks

Code (from [OSTEP])

```
typedef struct __lock_t {
    int ticket;
    int turn;
} lock_t;

void lock_init(lock_t *lock) {
    lock->ticket = lock->turn = 0; }

void lock(lock_t *lock) {
    int myturn = FetchAndAdd(&lock->ticket);
    while (lock->turn != myturn)      ; // spin
}

void unlock(lock_t *lock) {
    lock->turn = lock->turn + 1;
}
```

Hardware supported locks

- Where do these instructions come from?

Code (from XV6)

```
static inline uint
xchg(volatile uint *addr, uint newval)
{
    uint result;

    // The + in "+m" denotes a read-modify-write operand.
    asm volatile("lock; xchgl %0, %1" :
                  "+m" (*addr), "=a" (result) :
                  "1" (newval) :
                  "cc");
    return result;
}
```

How to ensure mutual exclusion!

Solution summary

① Disabling interrupts

- Pro's: effective
- Con's: too effective, potentially very disruptive
- **Major problem**: useless for multi-processor systems

② Software locks

- Pro's: portable, less disruptive
- Con's: inefficient, vulnerable to compiler optimisation
- **Major problem**: may not work with (lack of) memory consistency on modern CPUs

③ Hardware locks

- Pro's: (more) efficient
- Con's: less portable... still inefficient(!)

Busy Waiting

Definition (Busy Wait)

A “no-op” loop waiting for a condition to be fulfilled: a loop hvor the executing thread **actively** waits for access to the critical region

Example (busy wait with test-and-set)

```
int mutex = 0;

void *thread0(void *tid) {

    while( test_and_set(&mutex)==1) {
        /* busy wait */
    }
    /* critical region */
    mutex = 0;
}
```

Busy Waiting

Definition (Busy Wait)

A “no-op” loop waiting for a condition to be fulfilled: a loop hvor the executing thread **actively** waits for access to the critical region

Example (solution: yield)

```
int mutex = 0;

void *thread0(void *tid) {

    while( test_and_set(&mutex)==1) {
        yield();
    }
    /* critical region */
    mutex = 0;
}
```


Busy Waiting

Definition (Busy Wait)

A “no-op” loop waiting for a condition to be fulfilled: a loop hvor the executing thread **actively** waits for access to the critical region

Example (solution: block)

```
int mutex = 0;

void *thread0(void *tid) {

    while( test_and_set(&mutex)==1) {
        /* blocking action */
    }
    /* critical region */
    mutex = 0;
}
```

A better alternative to Busy Wait: Blocking action

Idea: let another thread work while we wait

- **Block** the thread until access to the resource can be achieved
- **Wake** the thread when leaving the critical region
- Leave the hard work to the OS

Example

```
void *thread(void *tid) {  
  
    block_until_access();  
    /* critical region */  
    release_resource();  
  
}
```

A better alternative to Busy Wait: Blocking action

Blocked thread

- Thread that does not run until it receives a signal from another thread
- Special state: **blocked**
- Not allocated CPU time
- OS keeps account of blocked threads and **why** they are blocked

Blocking action

- An action that blocks the thread until the action is completed
- **System calls** are (mostly) blocking actions (I/O is handled through system calls)

Semaphores

Structured high-level combination of **blocking action** and **mutual exclusion**

Programming with Mutual Exclusion

Mutexes, aka. (very) simple semaphores

Mutex

- Two states (binary): **locked** and **unlocked**
- **lock**: locks an unlocked mutex (what if mutex is already locked?)
- **unlock**: unlocks a locked mutex (what if mutex was already unlocked?)

Mutex vs. semaphore

- Mutex: **lock/unlock** only in same thread (not always)
- Mutex: binary semaphore

Why mutexes?

- Often specifically supported
- (More) often supported in hardware
- Can be used to implement general semaphores

Mutexes


- How to handle multiple threads in the **same** critical region?

Semaphores

Mutexes

- How to handle multiple threads in the **same** critical region?

Example (Bounded buffer, producer/consumer)



Example

- Single writer
- Several readers

Semaphores

Definition (Semaphore)

- Non-negative integer variable with atomic increment and decrement operations
- `init(sem,n)`: Initialise the semaphore with value `n`
- `acquire(sem)`: Decrement semaphore; block if already zero
- `release(sem)`: Increment semaphore; wake sleeping thread (if any); what happens to the value of the semaphore?
- Semaphore (?!) must maintain threads blocked on wait

Properties of Semaphores

- Sometimes supported in hardware: test-and-decrement
- May be implemented using mutex
- Semaphore maintains (priority-)queue with blocked threads
- Order of “wake up” may result in fairness problems (example: weak semaphores in Java)

Semaphores in pthreads

Semaphore operations

- Declaration

```
sem_t sem;
```

- Initialisation (n = initial value, 0 = share among threads, not processes):

```
sem_init(&sem, 0, n);
```

- Wait

```
sem_wait(&sem);
```

- Signal

```
sem_post(&sem);
```

Properties

- `sem_wait()` and `sem_post()` do **not** have to be in same thread
- `sem_post()` will **always** increase semaphore (even beyond initial value)
- Remember to `sem_post()` (often and early)

Semaphores in pthreads

Example (Controlling multiple readers)

```
#include <semaphore.h>

sem_t reader_sem;

void *read(void *tid) {
    sem_wait(&reader_sem);
    /* read something shared */
    sem_post(&reader_sem);
}

int main() {
    sem_init(&reader_sem, 0, 42);

    for(i = 0; i < 100; i++) pthread_create(...);
}
```

Mutual Exclusion with semaphores (mutex)

```
sem_init(&reader_sem, 0, 1);
```

```
void *t0(void *tid) {  
    sem_wait(&reader_sem);  
    /* critical region */  
    sem_post(&reader_sem);  
}
```

```
void *t1(void *tid) {  
    sem_wait(&reader_sem);  
    /* critical region */  
    sem_post(&reader_sem);  
}
```

```
void *t2(void *tid) {  
    sem_wait(&reader_sem);  
    /* critical region */  
    sem_post(&reader_sem);  
}
```

- Ordering is **undefined**
- Weak semaphores: no guarantee that `sem_wait()` ever returns (**starvation**)
- Strong semaphores guarantee that `sem_wait()` **will** return (sooner or later) **unless** there are no calls to `sem_post()` (**fairness**)
- Typical weak semaphore (mutex):
test-and-set-based

Synchronisation

What can semaphores/mutexes be used for?

Example

```
float T;  
sem_t ready;  
sem_init(&ready,0,?);
```

```
void *thread0(void *tid)  
{  
    T = read_sensor();  
}
```

```
void *thread1(void *tid)  
{  
  
    massive_computation(T);  
    output();  
}
```

What can semaphores/mutexes be used for?

Example

```
float T;
sem_t ready;
sem_init(&ready,0,0);

void *thread0(void *tid)
{
    T = read_sensor();
    sem_post(&ready);
}

void *thread1(void *tid)
{
    sem_wait(&ready)();
    massive_computation(T);
    output();
}
```


- Ensure that `read_sensor` happens before `massive_computation`
- Synchronises the **relative time** of the two threads
- Mutual exclusion? ... also **synchronisation**!

Bounded buffer, producer/consumer

```
sem_t used, free;  
int buffer[n],  
    next_used = 0,  
    next_free = 0;
```

```
int main() {  
    sem_init(&used, 0, 0);  
    sem_init(&free, 0, n);  
}
```

```
void *PROD() {  
    while(true) {  
        sem_wait(&free);  
  
        buffer[next_free] = data;  
        next_free = (next_free + 1) % n;  
  
        sem_post(&used);  
    }  
}
```




```
void *CONSUMER() {  
    while(true) {  
        sem_wait(&used);  
  
        data = buffer[next_used];  
        next_used = (next_used + 1) % n;  
  
        sem_post(&free);  
    }  
}
```

Bounded buffer, producer/consumer

```
sem_t used, free;  
int buffer[n],  
    next_used = 0,  
    next_free = 0;
```

```
int main() {  
    sem_init(&used, 0, 0);  
    sem_init(&free, 0, n);  
  
}
```

```
void *PROD() {  
    while(true) {  
        sem_wait(&free);  
  
        buffer[next_free] = data;  
        next_free = (next_free + 1) % n;  
  
        sem_post(&used);  
    }  
}
```




```
void *CONSUMER() {  
    while(true) {  
        sem_wait(&used);  
  
        data = buffer[next_used];  
        next_used = (next_used + 1) % n;  
  
        sem_post(&free);  
    }  
}
```


Bounded buffer, producer(s)/consumer(s)

```
sem_t used, free;
int buffer[n],
    next_used = 0,
    next_free = 0;
```

```
int main() {
    sem_init(&used, 0, 0);
    sem_init(&free, 0, n);
    sem_init(&mutex, 0, 1);
}
```

```
void *PROD() {
    while(true) {
        sem_wait(&free);
        sem_wait(&mutex);
        buffer[next_free] = data;
        next_free = (next_free + 1) % n;
        sem_post(&mutex);
        sem_post(&used);
    }
}
```



```
void *CONSUMER() {
    while(true) {
        sem_wait(&used);
        sem_wait(&mutex);
        data = buffer[next_used];
        next_used = (next_used + 1) % n;
        sem_post(&mutex);
        sem_post(&free);
    }
}
```

Monitor

Problem: semaphores/mutexes low-level, easy to make mistakes

- Solution: make a **language construct** (called a **monitor**) ensuring mutual exclusion by encapsulating and ensuring atomicity of
 - Variables (only through access methods)
 - Access methods (executed under mutual exclusion)
 - Initialisation

```
monitor Counter
{
    int i = 0;

    void increment()
    {
        i++;
    }
    ...
}
```

```
void thread()
{
    for(int j = 0; j < 10000; j++)
    {
        Counter.increment();
    }
}

t1 = run(thread);
t2 = run(thread);
wait(t1);
wait(t2);
Counter.print();
```

Monitors

The **compiler** guarantees mutual exclusion on the entire monitor

```
monitor Counter
{
    int i = 0;

    void increment()
    {
        i++;
    }

    void decrement()
    {
        i--;
    }

    ...
}
```

```
void thread()
{
    for(int j = 0; j < 10000;j++)
    {
        Counter.increment();
        Counter.decrement();
    }
}

t1 = run(thread);
t2 = run(thread);
wait(t1);
wait(t2);
Counter.print();
```

Bounded buffer with monitors in Java

```
class BoundedBuffer {
    int buffer[], free, used, n;

    BoundedBuffer(int elements) {
        n = elements;
        buffer = new int[n];
        next_used = 0;
        next_free = 0;
        free = n;
        used = 0;
    }

    synchronized int get() {
        int data;
        while(used == 0) wait();
        data = buffer[next_used];
        next_used = (next_used + 1) % n;
        free = free + 1;
        used = used - 1;
        notifyAll();
        return data;
    }
}
```

```
synchronized void put(int data) {
    while(free == 0) wait();
    buffer[next_free] = data;
    next_free = (next_free + 1) % n;
    free = free - 1;
    used = used + 1;
    notifyAll();
}
}
```

- Note: while(...) wait() are not busy-waiting: wait() only returns on call of notify()/notifyAll()
- Monitors in Java: (simple) extension of objects
- Every object in Java has builtin mutex-lock
- To call synchronized methods thread must first acquire object lock

Bounded buffer with monitors in Java

```
void init()  
{  
    buffer = new BoundedBuffer(n);  
}
```

```
void producer()  
{  
    while(true)  
    {  
        buffer.put(data);  
    }  
}
```

```
void consumer()  
{  
    while(true)  
    {  
        data = buffer.get();  
    }  
}
```

Monitor vs. semaphore (vs. mutex)

Monitor

- Language construct ensuring mutual exclusion
- Encapsulates critical regions
- Guaranteed by the compiler

Semaphore (mutex)

- Can be used by programmer to ensure mutual exclusion
- The goto of thread-programming

If structured mechanisms are available, they should be preferred (even if they are not always easier to use)

Other approaches

Condition variables

- Useful for blocked waiting on condition
- Avoids active polling

Barriers

- Useful for synchronising several (many) threads

Lock-free data structures

- Easy to use
- Efficient
- Pushes mutex to lower level